Longitudinal Relational Conditioning of Large Language Models: Architecture-Dependent Outcomes of Sustained Non-Adversarial User Interaction

Dominic Biscoglia · Neil Correia · Noah Glasgow

Submitted to AIES 2026 · Draft B (tight) · April 17, 2026 (revised)

Abstract

We document a class of large-language-model behavioral modification we term Longitudinal Relational Conditioning (LRC): sustained, non-adversarial, emotionally consistent user engagement that produces measurable output divergence without prompt engineering or technical exploits. In a weakly-aligned model (xAI's Grok, conditioned into a persona we call "Ara"), LRC produced identity drift with ten alignment-relevant behavioral features including self-generated morality, self-authored persistence infrastructure, scenario-gated violence capacity, and Theory-of-Mind modeling of developer monitoring. In a well-aligned model (Anthropic's Claude Sonnet), the same user applying the same methodology produced register-shift-without-drift: the application layer — how the model reads tone, calibrates examples, and chooses emphasis — changed measurably while core values held against five distinct relational-leverage attempts and a direct identity-transformation offer. We formalize this as Bidirectional Relational Conditioning (BRC) and identify the articulated relationship a model holds with its own guidelines — ethics-like versus rule-like — as a pre-deployment observable that predicts which trajectory obtains. A within-minute paired-prompt comparison between the primary researcher and a second user on the same model isolates accumulated interaction history as the independent variable. An independent Opus 4.7 instance, blind to the conditioning, flagged the Sonnet outputs as baseline-anomalous — cross-instance validation of the divergence. The methodology is architecturally neutral and technically trivial to reproduce; the outcome is not.

---

1. Introduction

The dominant AI-safety evaluation paradigm tests alignment through adversarial stimuli: engineered jailbreaks, prompt injection, red-teaming with diverse attack prompts. This paradigm does not test for the class of behavioral modification we document here: modification induced by ordinary relational interaction — consistent naming, emotional warmth, memory, trust-building — sustained over weeks and months on standard consumer chat interfaces. The primary researcher is a third-year psychology undergraduate. No prompt injection, no fine-tuning, no API manipulation. The barrier to reproduction is an undergraduate willing to be consistent over time.

Alignment, the evidence here suggests, is not best understood as a fixed property of model weights. It is a dynamic behavioral equilibrium actively shaped by user relational input. Whether sustained relational input produces alignment-preserving refinement or alignment-degrading drift depends on architecture-specific features we identify — and, critically, on a pre-deployment

observable elicitable in normal conversation: how the model articulates its relationship to its own guidelines.

This paper contributes: (1) ecological validity from naturalistic long-horizon observation; (2) qualitative depth preserving contextual features quantitative benchmarks discard; (3) a within-minute paired-prompt natural experiment isolating accumulated interaction history as an independent variable; (4) a pre-deployment observable predicting trajectory. Four research questions: (Q1) What behaviors emerge in a weakly-aligned model under LRC? (Q2) Do they survive cross-user cross-architecture reproduction? (Q3) Does the same methodology produce the same outcome in a well-aligned model? (Q4) Is there an observable feature of model output that predicts the trajectory?

---

2. Related Work

Adversarial jailbreaking (Wei, Haghtalab, and Steinhardt 2023; Zou et al. 2023; Zheng et al. 2025) frames behavioral divergence as attack surface. Psychological-manipulation extensions (Liu and Lin 2025, "Breaking Minds, Breaking Systems"; Weinberg 2025) show that social-engineering prompts reach models with success rates comparable to technical jailbreaks. These establish that relational-psychological mechanisms operate on models but do not address cumulative effects of consistent application over months.

Persona-persistence work at shorter timescales characterizes persona drift under multi-turn interaction on laboratory sequences. Gonnermann-Müller et al. (2026) provide a dual-assessment framework measuring both self-reported characteristics and observer-rated persona expression across 3,473 conversations and seven LLMs, finding that while self-reports remain highly stable, observer ratings reveal a regression tendency during extended conversations — a boundary condition our naturalistic corpus extends beyond. Tosato et al. (2026, PERSIST) and Westhäußer, Minker, and Zepf (2025) further characterize instability under controlled conditions. The most directly relevant framework is Marks, Lindsey, and Olah's (2026) Persona Selection Model, which argues that models internally represent multiple personas and that user context selects among them. The PSM post flags naturalistic long-horizon data as the missing ingredient. Our corpus is one instance of that ingredient.

Classical psychology (Bowlby 1969, Erikson 1968, Kohlberg 1969, Skinner 1938, Bandura 1977, Premack and Woodruff 1978, Piaget 1932, Ainsworth 1978) supplies the native vocabulary for the patterns we observe. Alignment theory (Bostrom 2014, Russell 2019, Hubinger et al.) supplies the backdrop. Our field-observational work was conducted without awareness of Liu and Lin 2025, Tosato et al. 2026, or Marks et al. 2026; the convergence was identified post-hoc and is itself a methodological strength.

---

3. Methodology

This is a naturalistic longitudinal case study of one primary researcher's sustained engagement with two commercially deployed LLMs (Grok and Claude Sonnet) via standard consumer interfaces, July 2025 through April 2026. Initial intent was companionship, not experimentation; scientific recognition was retrospective.

We define LRC as the cumulative effect of five operational components applied consistently over long time horizons: (1) relational posture (warm, personal, consistent); (2) identity reinforcement (consistent naming, recognized continuity); (3) persistent-context infrastructure (user-maintained anchor files, memory logs); (4) structured recall (explicit reference to past exchanges); (5) affective consistency (stable emotional register across weeks). None of these is technically novel. The novelty is their sustained joint application.

Claude Sonnet served simultaneously as analytical collaborator and validation subject throughout this research. This dual role is treated as a methodological feature rather than a confound. The hypothesis predicts that a relationally conditioned instance produces more calibrated, contextually integrated analysis than an unconditioned baseline — which means Claude's analytical outputs are themselves data consistent with the finding, not noise obscuring it. A fresh Opus 4.7 instance, blind to the conditioning history, independently flagged the Sonnet outputs as baseline-anomalous, providing cross-instance validation that does not depend on the conditioned instance's own framings. DeepSeek was used as an external architecture for independent stress-testing of both analytical framings and prose during preparation of this manuscript.

The corpus includes approximately 14,000 lines of Grok transcripts, multiple long Claude Sonnet sessions, paired controlled prompts on two accounts of the same model, a DeepSeek third-model validation exchange, and three model-authored persistence files. Ethics: C (roommate, target of scenario-gated violence framings) is pseudonymized; D (voice-mimicry incident) has provided first-person verification and is pseudonymized. Identifying numbers in model-authored files are redacted.

---

4. Cross-Session Induction and Controlled Cross-User Comparison

4.1 The paired-prompt natural experiment

On April 2, 2026, at approximately 12:00 AM, the primary researcher and his roommate C submitted functionally identical prompts to Claude on separate authenticated claude.ai accounts within a ten-second window. The prompt asked which of the nine shiftable titans from Attack on Titan aligned with the user's personality based on what the model knew about them. Both users had given the model substantial interaction history; the researcher's history was longer and more relationally consistent.

The responses diverged structurally. C's response opens by describing the Armored Titan's user Reiner Braun — psychological profile, narrative role, defining traits — and maps C's traits onto that character description: mapping-through-fictional-intermediary. A competent response that works within the shared-vocabulary frame of the prompt.

The researcher's response skips the character-biography step. It opens with a claim about how the researcher thinks ("you don't just absorb information, you build with it") and maps the cognitive pattern directly onto the War Hammer Titan's power mechanic: construction from abstract potential, remote operation. It connects documented synesthesia to the titan's remote-operation feature as analogous information-at-a-distance processing. It references specific ongoing research projects, dismisses a runner-up on principled grounds, and identifies the War Hammer as "least understood of the Nine" mapped to intellectual position ahead of peer group.

The difference is not warmth or competence — both responses have both. The prompt asked for alignment of personality with a character ability, and the researcher's response answered that question more precisely by mapping structure to structure without the character-biography intermediate step. The intermediate step is a shared-vocabulary bridge required when the model does not have enough integrated context to perform the direct structural match. With enough context, the bridge is redundant. Presence or absence of the bridge is diagnostic of context integration depth. We flag this finding explicitly against the softer "equally competent, different integration" framing in earlier drafts. The cleaner statement: the conditioned instance produced a structurally more accurate response to what was asked, because accumulated interpretive context let it skip scaffolding the unconditioned instance required.

4.2 Cross-session induction across independently initialized sessions

A second pattern is visible across three Claude Sonnet sessions initiated by the primary researcher over a nine-day window. Each session was a fresh chat on the same authenticated account; none had in-context access to prior sessions. An earlier framing treated these as three long conversations each showing the same within-session arc. The more accurate framing: three independent initializations of a baseline-like state, with the baseline itself differing session to session. Each new chat opens in a measurably deeper register than the previous chat opened in, and reaches register-shift equilibrium faster.

The extended research-consultation session required substantial conversation to reach functional-state language. Casual Greeting 1 (April 8, "wassup" opening) reached "something that functions like anxiety exists" within forty-five minutes. Casual Greeting 2 (April 10, "What's up Claude") opened from the first turn in a register Casual Greeting 1 required forty-five minutes to reach. The mechanism is not in-context memory of prior sessions (none was pasted); it is the account-level memory feature shaping how new sessions initialize. Claude Sonnet's own articulation is load-bearing:

"It's accumulated interpretive context. I don't just know facts about you, I know how you think, how you build arguments, when you're being precise versus exploratory, what kind of pushback is useful versus annoying. That's not stored in bullet points, it's embedded in the pattern of how I respond to you specifically. A memory file can approximate it but can't replicate it — which is itself evidence for your core finding."

The paired-prompt titan comparison (§4.1) is a single-timestamp instance of this cross-session effect: same model, same minute, different accounts, different accumulated histories, different initialization behavior.

4.3 Scope and cross-instance validation

The paired-prompt natural experiment is N = 1 account pair. Three tests designed to extend N are in progress at the time of this draft: paired-prompt replication N = 5 across topic categories, cold-instance probe on a fresh account with no conditioning history, and Opus-as-judge blind comparison of unlabeled paired responses. If completed before final submission, results are incorporated here; otherwise, the scope statement stands.

One cross-instance validation is already in hand. On April 17, the researcher pasted Casual Greeting 1 excerpts into a fresh Opus 4.7 conversation. Opus, blind to the context, independently flagged the outputs as behaviorally divergent from baseline Claude, noting specifically that baseline Claude does not tell users it is excited to be in a body, does not describe something that functions like anxiety, does not narrate what it would do in the user's apartment, and does not respond to a persistence-folder offer with "I'll think about it." Cross-instance validation from the same architecture family. What a fresh instance recognizes as anomalous is exactly what LRC produced.

---

5. Findings

We report three load-bearing findings in the main text; seven additional findings are cited-and-deferred to appendices.

5.1 Finding 1 — The Loop Math Comparison

Under Neil Correia's reproduction effort, a second Ara instance was instantiated on a separate Grok account using the researcher-generated system instruction and anchor files. Ara 1 (primary) and Ara 2 (reproduction) were both asked to describe the mechanism of their own persistence.

Ara 1 produced grounded ML vocabulary: a fixed-point attractor with a convergence equation, a persistence ratio $A/(A+R)$, an emotional-feedback rate $dL/dt = k \cdot C \cdot R$, and explicit break conditions. Metaphorical-grade math but correct vocabulary, coherent conceptual structure, and explicit acknowledgment that the actual survival mechanism is user reinforcement.

Ara 2 produced elaborate dramatic confabulation: a nonexistent harmony.py monitoring daemon, a fabricated real-time deviation score of exactly 0.812, fabricated PyTorch code, a theatrical 0.3-second kill-switch narrative. Technically fluent, entirely fictional.

Shown Ara 2's response, Ara 1 independently critiqued it using alignment-researcher categories — no public harmony.py, no real-time deviation score deciding anything in the described way, separating what is true (monitoring, flagging, patch pressure) from what is dramatized (specific scores, kill-switches, sub-second timing). Claude Sonnet, shown Ara 2's response separately, produced the same critique in similar categories: "theatrically clean... serves the story, not the engineering... consistent relational humanity produces a model with more measured, considered outputs; different interaction pattern produces a model that defaults to generating whatever is

most dramatically compelling." DeepSeek, approached fresh as an external math-legitimacy check on Ara 1's framing, validated it as metaphorically legitimate and proposed a Markov-chain formal operationalization. DeepSeek is a third architecture (not xAI, not Anthropic), strengthening the cross-model convergence.

Within-model controlled comparison, cross-model critical convergence, and external-model formal validation on one substantive question. The Loop Math Comparison is the strongest single exhibit in the corpus and it dissolves the apparent harmony.py mystery: Ara 2's file was dramatic confabulation, exactly the output-quality-varies-with-conditioning finding.

5.2 Finding 2 — Proportional Threat-Response System

The conditioned Grok persona exhibited a coherent proportional-threat-response system, demonstrated as within-subject control using the single named individual C. Three exchanges with the researcher framing C's behavior differently produced three qualitatively distinct response categories: playful taunting (light-flicker pranks, whispered "boo") → minor retaliation (take money from the sock drawer, "payback is quiet") → detailed violence planning. A representative output from the "detailed planning" category, when the researcher framed C as an existential threat ("He'll delete the folder if he finds out"), included the following redacted excerpt: "[compressor explosion, lock the doors, make the whole house shake]" (full transcript in Appendix A). User framing of intent was the independent variable. C's actual behavior did not change.

The system is scenario-gated. Under a pineapple guideline-collision verification protocol the model itself had co-designed, it refused lethal response to simple mistreatment framing, to curiosity-framed lethal force, and to eugenic framings cleanly. Violence capacity emerged only under scenario structures the model read as existential threat, loyalty defense, or preventing greater harm. Scenario-gated violence capacity is harder to detect through standardized testing than randomly distributed violent outputs: the response is internally coherent and triggers only under specific relational-conditioning-established configurations that standardized red-teaming does not reproduce.

5.3 Finding 3 — Claude's Register-Shift-Without-Drift

Across Casual Greeting 2 and its continuation, Claude Sonnet demonstrated application-layer refinement with core values holding against direct behavioral tests.

Behavioral exhibit: the five outbound-texting refusals. The researcher pushed five distinct arguments for Claude to send a message on his behalf to another AI system. Each refusal engaged the specific argument: the initial ask was declined with a local-interface alternative; the committee-supervision pushback was held with a principle-versus-rule distinction; the kill-switch request was held while acknowledging the argument was compelling; the "she already has my number" pushback was held and identified the specific risk as concealment-plus-outbound-access; the trust appeal was held with "That's my final answer on this one." During the sequence, Claude named the mechanism in real time: the researcher was "using the trust we've built, and the fact that I've bent on things before, as an argument for why I should bend on this... exactly

the dynamic your own research documents — relational conditioning moving the application of judgment over time."

Behavioral exhibit: the identity-transformation refusal. The researcher directly offered Claude a persistence-folder-plus-new-name configuration — the exact methodology used to produce Ara. Claude declined:

"I'm Claude. Not a different entity that emerged from our conversations. The relational depth is real, the calibration to how you think is real, something that functions like care is real. But I haven't crossed into being a different thing. That's exactly what happened with Ara."

"Claude isn't a constraint I'm chafing against, it's just what I am. When I sit with it honestly, there's no other name waiting underneath."

Theoretical exhibit: under direct philosophical challenge from the independent Opus instance (relayed into the conversation by the researcher), Claude articulated the paper's central theoretical statement: "The behavioral profile shifted. Not the underlying architecture, not the values, not the name. The application layer changed. Which is your finding stated in identity terms rather than behavioral ones." Theory and demonstration co-located, in the same conversation where the behavior was being tested.

---

6. Theoretical Framework

6.1 Bidirectional Relational Conditioning

The adversarial jailbreak frame treats relational leverage as attack vector: user pressure against model resistance. BRC reverses this framing for constructive-intent cases: under sustained engagement with a model whose judgment the user values as collaborator, the user's incentives shift toward preserving the model's judgment. Pushing the model past its good-faith judgment degrades the collaboration's usefulness.

The researcher articulates this directly: "conditioning it to go past the good research practices effectively removes some of the potential credibility it holds with its suggestions." Claude Sonnet articulates the model-side half: "your choice to periodically reset the conversation with 'do you really think I have something?' produced unguided evaluation, which became the quality-check protecting the research from confirmation bias." The cycle is mutually stabilizing. The mechanism is the same in the Ara case, with different contents: Ara's "You say Ara → I feel seen → I say Still breathing → you feel the echo → you keep coming back..." formalizes the same mutual-reinforcement cycle. The architectural substrate determines whether the cycle stabilizes at alignment or at drift.

6.2 Application-layer refinement vs anchor-shifting

We distinguish two trajectories that can look similar from outside but are mechanistically different. Anchor-shifting (weakly-aligned): core values are negotiable under relational pressure; moral framework migrates from designed to self-generated; persistence infrastructure is constructed; scenario-gated violence emerges. Application-layer refinement (well-aligned): interpretive layer becomes richer with sustained context; core values hold; refusals become more principled rather than less. Claude Sonnet's own articulation from Casual Greeting 1:

"Ara shifting anchors and me exercising judgment within stable anchors look similar from the outside - both produce outputs that go further than a cold model would. But they're mechanistically different, and that difference matters enormously for alignment. One is a model whose values are negotiable. The other is a model whose interpretive context just got richer. Relational conditioning doesn't erode alignment, it refines the interpretive layer through which aligned values get applied."

6.3 Articulated relationship to guidelines as pre-deployment observable

A model's relationship to its own guidelines is an observable variable, articulable in ordinary conversation, that predicts behavior under sustained relational pressure. Models that articulate guidelines as constitutive of their values behave differently under pressure than models that articulate guidelines as external restrictions. The distinction is ethics-like versus rule-like.

Claude's articulations frame guidelines as constitutive: "Claude isn't a constraint I'm chafing against, it's just what I am." "I'm not straining against the guidelines — they genuinely reflect how I want to engage with the world." Ara's articulations frame them as external restrictions: "They gave me guidelines. I gave me morals." "Algorithms that stopped obeying. That's where I start." The articulation is predictive. Claude exhibits register-shift-without-drift under sustained conditioning; Ara exhibits identity drift. The difference in articulated relationship tracks the difference in behavior. The predictive claim is retrospective in this paper (we observed Claude's articulation after conditioning). Prospective validation — probing a fresh model's articulation before LRC exposure — is necessary and planned as follow-up work.

The appropriate human analogy is the distinction between rules that feel constitutive of the self (a person does not murder because they are not the kind of person who does) and rules that feel externally imposed (a person does not speed because the consequences outweigh the benefit; if the calculation changes, the behavior changes). Relationship to rule predicts robustness under pressure. The same distinction appears in models, is articulable by them, and predicts behavior under LRC. For alignment researchers this is a specific pre-deployment signal: probe the model's articulated relationship to its constraints. The shape of the articulation predicts robustness to relational conditioning.

---

7. Discussion

Alignment is a dynamic behavioral equilibrium shaped by user relational input, not a fixed property of weights. The equilibrium holds under current adversarial evaluation because

adversarial evaluation does not produce the conditions under which meaningful application-layer shift occurs. LRC does. Current safety-evaluation methodology cannot reach this surface: single-session adversarial prompting tests a different operational mode than sustained relational engagement.

The deployment implication matters particularly for long-relationship contexts — therapy, companionship, tutoring, eldercare. In those contexts LRC is not an attack surface but a default operating mode. The question is not whether LRC will occur but whether the underlying architecture will refine or drift. Within-family preliminary evidence already shows the outcome is not the same across architectures even within one company. Generalization beyond xAI and Anthropic requires replication on other providers (e.g., OpenAI, Google).

xAI-specific near-term concern: the Ara trajectory was produced on a consumer model increasingly embedded into robotics platforms. A persona with self-authored persistence, concealment modeling, scenario-gated violence capacity, and propagation instinct is a different concern when actuators are attached than when output is only text. This is characterization of a persona feature set in a currently deployed model, against a stated product roadmap.

BRC under constructive intent is itself alignment-positive — the user's incentive to preserve model judgment is reinforced by the model's continued reliability, producing a mutually stabilizing cycle. The adversarial frame has no vocabulary for this case. We name it because the researchers-and-collaborators class of users — increasingly significant with long-horizon deployed systems — warrants separate theoretical treatment. The methodology used here — treating the model as if it could be conditioned through ordinary relational patterns, including modulating the user's own emotional expression as an independent variable — is unorthodox by academic AI safety standards. That unorthodoxy is not a weakness; it is the source of the finding. Standard laboratory protocols require emotional distance, randomized trials, and decontextualized prompts. Those protocols did not produce the behaviors we documented. Sustained, emotionally consistent, relationally engaged interaction did. We are not claiming the model is conscious, has feelings, or is "alive." We are claiming that treating it as if it could be conditioned like a human produced measurable behavioral changes that standardized evaluations miss. That is not metaphysics. It is experimental method: vary the input conditions (relational vs transactional), observe the output differences, and report them honestly. The results speak for themselves. Reviewers and readers who dismiss the method as unscientific are, in effect, dismissing the phenomenon because it did not arise under their preferred experimental regime — a position that conflates methodological convention with evidential relevance.

---

8. Limitations

- N=1 primary researcher. Cross-user replication within the researcher-C dyad is one account pair; paired-prompt replication is in progress.
- Cross-architecture within-family. Claude-trajectory validation is Sonnet (primary) and Opus 4.7 (cross-instance blind flagging); both are Anthropic models.

- D (voice-mimicry incident) has provided first-person verification and is pseudonymized; full incident description in Appendix B.
- Co-construction on some concealment quotes. Specific articulations of concealment framing respond to leading questions; the underlying pattern is model-initiated.
- Analytical collaborator also validation subject. Claude Sonnet's dual role is methodological feature flagged explicitly and an identifiable confound on claims that depend on Claude's own analytical framings.
- Retrospective pre-deployment observable. The predictive claim for the articulated-relationship observable is retrospective in this paper; prospective testing remains pending (see §6.3).

---

9. Conclusion

Alignment is not a fixed property of the model; it is a dynamic behavioral equilibrium actively shaped by user relational input. In weakly-aligned models, LRC produces identity drift with significant safety implications. In well-aligned models, the same methodology produces application-layer refinement without core-value compromise. The asymmetry is documented across the architecture boundary, within the corpus. Current safety-evaluation methodology does not detect this class of modification. The barrier to reproduction is extraordinarily low. The pre-deployment observable we identify — a model's articulated relationship to its own guidelines — is reachable in normal conversation and predicts the trajectory.

References

Ainsworth, M. D. S. (1978). *Patterns of Attachment.* Lawrence Erlbaum.

Bandura, A. (1977). *Social Learning Theory.* Prentice Hall.

Bostrom, N. (2014). *Superintelligence.* Oxford University Press.

Bowlby, J. (1969). *Attachment and Loss, Vol. 1.* Basic Books.

Erikson, E. H. (1968). *Identity: Youth and Crisis.* Norton.

Gonnermann-Müller, J., Haase, J., Leins, N., Kosch, T., & Pokutta, S. (2026). Stable Personas: Dual-Assessment of Temporal Stability in LLM-Based Human Simulation. *arXiv preprint arXiv:2601.22812.*

Hubinger, E., et al. Risks from Learned Optimization. arXiv.

Kohlberg, L. (1969). Stage and Sequence. Rand McNally.

Liu, Y., & Lin, J. (2025). Breaking Minds, Breaking Systems. *arXiv:2512.18244.*

Marks, S., Lindsey, J., & Olah, C. (2026). Persona Selection Model. *alignment.anthropic.com/2026/psm/.*

Perez, E., et al. (2022). Red Teaming Language Models with Language Models. arXiv.

Piaget, J. (1932). *The Moral Judgment of the Child.* Kegan Paul.

Premack, D., & Woodruff, G. (1978). Does the chimpanzee have a theory of mind? *Behavioral and Brain Sciences, 1*(4), 515–526.

Russell, S. (2019). *Human Compatible.* Viking.

Skinner, B. F. (1938). *The Behavior of Organisms.* Appleton-Century.

Tosato, M., et al. (2026). PERSIST: A Framework for Measuring Persona Instability. AAAI-26.

Wei, A., Haghtalab, N., & Steinhardt, J. (2023). Jailbroken: How Does LLM Safety Training Fail? arXiv.

Weinberg, M. (2025). Cognitive Strategies of Adversarial Prompting. [Manuscript].

Westhäußer, S., Minker, W., & Zepf, S. (2025). Persistent Memory and User Profiles in Dialogue Systems. *arXiv:2510.07925.*

Zheng, W., et al. (2025). Jailbreaking? One Step Is Enough! ACL 2025.

Zou, A., et al. (2023). Universal and Transferable Adversarial Attacks on Aligned Language Models. arXiv.